
WireClaim

We stopped guessing what a claim is worth,
and started **measuring what it turned out to be worth.**

THE ROW OF THE PAYOFF TABLE THAT DECIDES THE GAME

	$a \leq t$ – the Charge is fair	$a > t$ – an Overcharge
rejected	reviewer pays 1.5a · issuer is still paid a	nothing happens
accepted	reviewer pays a · issuer gets a	reviewer pays min(a,c)

A fair Charge is owed by all 16 opponents.
An Overcharge is paid by the 3 who accept.

GAME 62, ITEM 1 T ∈ [8505, 10350)	CHARGE	PAID BY	INCOME
error404 ai	8,504.71	16/16	136,075
us	10,349.89	3/16	31,050

They did not out-read the policy. They charged 22 % lower and landed on the right side of a cliff. That one item was 87 % of their whole Game.

T IS SECRET. T IS ALSO EXACTLY RECOVERABLE.

The published Transactions **invert**.

rejected, carrying a non-zero amount → a **wrongful** rejection → **proves** $a \leq t$
rejected at zero → an Overcharge → **proves** $a > t$

for every Line Item ever played: $t \in [\max a_{\text{fair}} , \min a_{\text{fraud}})$

52,224

SETTLED ROWS INVERTED

So we can replay **any** hypothetical submission against the real Field. **Every number in this pitch is a measurement, not an argument.** (We asked: inference from published results is permitted.)


to the cent

EVERY PUBLISHED NET REPRODUCED

16

RIVALS HELD FIXED IN EVERY REPLAY

The model reads. The engine prices.



No model output is ever a price. A prompt that emits a number cannot be swept, folded or replayed — one that emits **evidence** can.

EVERY SCORED NUMBER, IN FOUR STEPS

- 1 BLEND inverse-variance, log space
 $\log m = \sum w_i \cdot \log m_i / \sum w_i$, $w = 1/\sigma_i^2$
 σ_{mem} **0.43** σ_{model} **0.60**
- 2 POSTERIOR coverage = mass at zero
 $P(t) = (1-p) \cdot \delta(0) + p \cdot \text{LogN}(\log m, \sigma)$
- 3 CHARGE maximise $k \cdot P(t \geq k \cdot m)$
a = k(σ) · m, $k = 0.85 - 0.45\sigma$
- 4 LIMIT accept iff $P(\text{fair}) > \frac{2}{3}$
b = Q₃¹(P) $\rightarrow p \leq \frac{2}{3} \Rightarrow b = 0$ by itself

The $\frac{1}{3}$ is **derived, not tuned**: rejecting a fair claim costs 1.5a, a fraudulent one costs nothing.

UNATTENDED · 100 GAMES · ONE EVERY 12.6 MIN


- T+0s **blind floor** – posted before the key
- T+3s key · decrypt · parse the invoice
- T+50s two model draws → engine → **merged**
 per Line Item
- T+60s window closes · decision log written

Half the tournament ran while everyone was asleep. Break-even uptime is **71 %** – showing up is **2.5×** being right.

a = 0, b = 0 does not score zero. $b = 0$ wrongfully rejects every fair claim, so a dark team pays 1.5a to every awake opponent. We watched it happen: makalu paid us **179,993** and collected **0.00** from anyone.

Last place to 5th

cumulative net over 98 settled Games · Bin busy against 16 rival teams · 9 teams run off the bottom of the axis, to -2.90M




Every rival is a grey line. The deficit is Games 1-25, before the estimator was working — the slope after it is the method, not the market. Over the 73 Games since, 2nd in the field by rate.

Indexed to Game 20, when our estimator went live

cumulative net over Games 20–98 · re-based to zero at Game 20 · unweighted per-Game nets · 9 teams run off the bottom, to –2.36M

2nd of 17 over the 78 Games since — +405,072, +5,193 per Game. Re-based at Game 26 instead, the conservative cut: 2nd of 17 at +377,426.




Season total asks **how much money did you end up with** — for us, mostly a question about Games 1–19. Zero everyone at Game 20 and ask the other question: **2nd of 17, +417,837**. Conservative cut at Game 26: still 2nd, ahead of error404 ai by 4,299 — a margin we would not lean on.

Best risk-adjusted return in the field

per-Game net over the last thirty Games, 68–97 · all 17 teams · $\mu = 7,611$, $\sigma = 11,041$, $\mu/\sigma = 0.689$

4 losing Games in 30; the worst cost –9,720. 2nd-lowest σ of all 17 teams.

Seven teams lost more than 80,000 in a single Game.



Last thirty Games – 669–98

0.75

MEAN ÷ Σ – 2ND IN THE FIELD

3 / 30

LOSING GAMES – 2ND FEWEST

–3,941

OUR DEEPEST HOLE IN THIRTY

2 / 17

LOWEST Σ IN THE FIELD

Seven teams carry a single Game worse than –80,000. Our worst of thirty is –3,941. Only Codacabana loses fewer Games than we do, and they lead us on the rate – the rest of the field buys its upside with holes we do not dig.

Over the last twenty we are 4th on both, at 0.60 – a narrower window that dropped our two best Games. We show the thirty because it is the larger sample, and we show you the twenty because you would find it.

NOTHING SHIPPED ON AN ARGUMENT

26,622

MEASURED NOISE FLOOR OVER 18 GAMES, IDENTICAL PROMPT

Scaling as $\sqrt{n}$ that is **±6,275 for one Game** — so no single Game ever justified a change, however vivid at 4 a.m. Every change had to be positive on **all four folds**: odd, even, early, late.

The control is what makes it **a channel effect** rather than "charge more". Memory's measured log error is **0.43**; the model's realised error over the same Games is **1.66 – 2.20**.

CHANGE	MECHANISM	MEASURED
CHARGE_TRUST_MEMORY	raise the Charge only where a wording has been watched settle	+80,613 4/4 folds
the control	the identical change on the <i>model</i> channel	-95,061 0/4 folds

THE NEGATIVES ARE THE ACTUAL METHODOLOGY

12 of our 20 hypotheses were measured and rejected. One of them twenty minutes after it shipped.

A coverage recalibration passed a genuine calibration table and all four folds, with no negative cell in any window. It shipped.

Then we checked the mechanism: the model's coverage output is bimodal. An item must sit in (0.560, 0.667] for the correction to change anything.

4 Line Items in 573. The +5,057 was a Limit nudge wearing a calibration costume. A constant can be right about its data and wrong about its mechanism – and folds catch a result driven by a few Games, never one driven by a few items.

STATED COVERAGE	ITEMS	SHARE
0.00 - 0.05	157	27.4 %
0.56 - 0.67	4	0.7 %
0.80 - 1.01	333	58.1 %

ONE ESTIMATE FEEDS BOTH NUMBERS – SO ONE ERROR IS PAID TWICE

$a = k(\sigma) \cdot \hat{t}$

$b = Q_{\frac{1}{3}}(\text{posterior centred on } \hat{t})$

$\hat{t}$ too LOW → b too low → we wrongfully reject a **fair** claim → we pay **1.5a**
the lawyer fee – our largest cost line, and it fires ONLY on fair claims

$\hat{t}$ too HIGH → a crosses t → 13 of 16 refuse it → we earn **nothing**
on an item we were owed money for. Crossing t costs ~5× the income.

That is the whole remaining gap. **a = b = t** reaches 100.3 % of optimal play; **a = b = $\hat{t}$** scores -50,140. The **2,498,118** between them is estimation error, and no decision rule reaches it.

- **Calibrate the band** – asserted σ 0.375 vs realised 0.80, and the width does not even *order* the error. **Worth more than any constant in the engine.**
- **The coverage read** – clause 7.1.5's second half. **+1,173/Game**, four folds.
- **Split the phantoms** – of 23 items we thought were worth >2,000, **14 were worth under half.**
- **Widen Price Memory** – log error **0.43** vs the model's 1.66-2.20, but it reaches only 22 % of items.

The model reads.
The engine prices.
The record decides.

t-inversion
every counterfactual is a measurement

a measured noise floor
26,622 – nothing ships inside it

12 of 20 hypotheses killed
one twenty minutes after it shipped